

Result

Ruiqi Wang

2024-05-22

```
library(keras)
mnist <- dataset_mnist()
train_images <- mnist$train$x
train_labels <- mnist$train$y
test_images <- mnist$test$x
test_labels <- mnist$test$y
str(train_images)
```

```
## int [1:60000, 1:28, 1:28] 0 0 0 0 0 0 0 0 0 0 0 ...
```

```
# int [1:60000, 1:28, 1:28] 0 0 0 0 0 0 0 0 0 0 0 ...
str(test_labels)
```

```
## int [1:10000(1d)] 7 2 1 0 4 1 4 9 5 9 ...
```

```
# int [1:10000(1d)] 7 2 1 0 4 1 4 9 5 9 ...
train_images <- array_reshape(train_images, c(60000, 28 * 28))
train_images <- train_images / 255
test_images <- array_reshape(test_images, c(10000, 28 * 28))
test_images <- test_images / 255
train_labels <- to_categorical(train_labels)
test_labels <- to_categorical(test_labels)
```

```
network <- keras_model_sequential() %>%
layer_dense(units = 512, activation = "relu", input_shape = c(28 * 28)) %>%
layer_dense(units = 10, activation = "softmax")
network %>% compile(
optimizer = "rmsprop",
loss = "categorical_crossentropy",
metrics = c("accuracy")
)
network %>% fit(train_images, train_labels, epochs = 5, batch_size = 128)
```

```
## Epoch 1/5
## 469/469 - 4s - loss: 0.3817 - accuracy: 0.8893 - 4s/epoch - 8ms/step
## Epoch 2/5
## 469/469 - 3s - loss: 0.3204 - accuracy: 0.9099 - 3s/epoch - 7ms/step
## Epoch 3/5
## 469/469 - 3s - loss: 0.3228 - accuracy: 0.9113 - 3s/epoch - 7ms/step
## Epoch 4/5
## 469/469 - 3s - loss: 0.3325 - accuracy: 0.9102 - 3s/epoch - 7ms/step
## Epoch 5/5
## 469/469 - 3s - loss: 0.3457 - accuracy: 0.9093 - 3s/epoch - 7ms/step
```

```
network_1 <- keras_model_sequential() %>%
layer_dense(units = 512, activation = "relu", input_shape = c(28 * 28)) %>%
layer_dense(units = 10, activation = "softmax")
network_1 %>% compile(
optimizer = "rmsprop",
loss = "categorical_crossentropy",
metrics = c("accuracy")
)
network_1 %>% fit(train_images, train_labels, epochs = 5, batch_size = 64)
```

```
## Epoch 1/5
## 938/938 - 7s - loss: 0.3785 - accuracy: 0.8926 - 7s/epoch - 8ms/step
## Epoch 2/5
## 938/938 - 7s - loss: 0.3450 - accuracy: 0.9064 - 7s/epoch - 7ms/step
## Epoch 3/5
## 938/938 - 7s - loss: 0.3658 - accuracy: 0.9050 - 7s/epoch - 7ms/step
## Epoch 4/5
## 938/938 - 7s - loss: 0.3950 - accuracy: 0.9020 - 7s/epoch - 7ms/step
## Epoch 5/5
## 938/938 - 7s - loss: 0.4291 - accuracy: 0.8990 - 7s/epoch - 7ms/step
```

Here the batch size decreases from 128 to 64. The time is longer than before. Also, the best accuracy is drop from 0.91 to 0.9.

```
metrics <- network %>% evaluate(test_images, test_labels, verbose = 0)
metrics
```

```
##      loss accuracy
## 0.351831 0.914200
```

```
network
```

```
## Model: "sequential"
##
## _____
## Layer (type)                Output Shape                Param #
## =====
## dense_1 (Dense)             (None, 512)                 401920
## dense (Dense)               (None, 10)                  5130
## =====
## Total params: 407050 (1.55 MB)
## Trainable params: 407050 (1.55 MB)
## Non-trainable params: 0 (0.00 Byte)
## _____
```

```
#Setting aside a validation set:
val_indices <- 1:10000
x_val <- train_images[val_indices,]
partial_x_train <- train_images[-val_indices,]
y_val <- train_labels[val_indices,]
partial_y_train <- train_labels[-val_indices,]
history <- network %>% fit(
  partial_x_train,
  partial_y_train,
  epochs = 20,
  batch_size = 512,
  validation_data = list(x_val, y_val)
)
```


```
## Epoch 1/20
## 98/98 - 1s - loss: 0.3146 - accuracy: 0.9152 - val_loss: 0.3220 - val_accuracy: 0.9135 - 938ms/epoch - 10ms/step
## Epoch 2/20
## 98/98 - 1s - loss: 0.3155 - accuracy: 0.9136 - val_loss: 0.3777 - val_accuracy: 0.8999 - 826ms/epoch - 8ms/step
## Epoch 3/20
## 98/98 - 1s - loss: 0.3217 - accuracy: 0.9118 - val_loss: 0.3634 - val_accuracy: 0.8994 - 832ms/epoch - 8ms/step
## Epoch 4/20
## 98/98 - 1s - loss: 0.3221 - accuracy: 0.9106 - val_loss: 0.3890 - val_accuracy: 0.8961 - 827ms/epoch - 8ms/step
## Epoch 5/20
## 98/98 - 1s - loss: 0.3281 - accuracy: 0.9120 - val_loss: 0.3414 - val_accuracy: 0.9122 - 822ms/epoch - 8ms/step
## Epoch 6/20
## 98/98 - 1s - loss: 0.3350 - accuracy: 0.9103 - val_loss: 0.3862 - val_accuracy: 0.9018 - 819ms/epoch - 8ms/step
## Epoch 7/20
```

```
## 98/98 - 1s - loss: 0.3314 - accuracy: 0.9107 - val_loss: 0.4745 - val_accuracy: 0.8785 - 820ms/epoch - 8ms/step
## Epoch 8/20
## 98/98 - 1s - loss: 0.3441 - accuracy: 0.9088 - val_loss: 0.3059 - val_accuracy: 0.9211 - 826ms/epoch - 8ms/step
## Epoch 9/20
## 98/98 - 1s - loss: 0.3451 - accuracy: 0.9099 - val_loss: 0.4170 - val_accuracy: 0.8864 - 831ms/epoch - 8ms/step
## Epoch 10/20
## 98/98 - 1s - loss: 0.3472 - accuracy: 0.9101 - val_loss: 0.5957 - val_accuracy: 0.8539 - 828ms/epoch - 8ms/step
## Epoch 11/20
## 98/98 - 1s - loss: 0.3494 - accuracy: 0.9081 - val_loss: 0.5696 - val_accuracy: 0.8660 - 831ms/epoch - 8ms/step
## Epoch 12/20
## 98/98 - 1s - loss: 0.3646 - accuracy: 0.9073 - val_loss: 0.4179 - val_accuracy: 0.8941 - 841ms/epoch - 9ms/step
## Epoch 13/20
## 98/98 - 1s - loss: 0.3659 - accuracy: 0.9064 - val_loss: 0.4436 - val_accuracy: 0.8894 - 830ms/epoch - 8ms/step
## Epoch 14/20
## 98/98 - 1s - loss: 0.3634 - accuracy: 0.9071 - val_loss: 0.3707 - val_accuracy: 0.9142 - 829ms/epoch - 8ms/step
## Epoch 15/20
## 98/98 - 1s - loss: 0.3691 - accuracy: 0.9073 - val_loss: 0.4129 - val_accuracy: 0.8933 - 823ms/epoch - 8ms/step
## Epoch 16/20
## 98/98 - 1s - loss: 0.3823 - accuracy: 0.9043 - val_loss: 0.4451 - val_accuracy: 0.8948 - 825ms/epoch - 8ms/step
## Epoch 17/20
## 98/98 - 1s - loss: 0.3822 - accuracy: 0.9062 - val_loss: 0.4687 - val_accuracy: 0.8909 - 825ms/epoch - 8ms/step
## Epoch 18/20
## 98/98 - 1s - loss: 0.3879 - accuracy: 0.9050 - val_loss: 0.3783 - val_accuracy: 0.9097 - 821ms/epoch - 8ms/step
## Epoch 19/20
## 98/98 - 1s - loss: 0.3992 - accuracy: 0.9031 - val_loss: 0.4457 - val_accuracy: 0.8979 - 825ms/epoch - 8ms/step
## Epoch 20/20
## 98/98 - 1s - loss: 0.3942 - accuracy: 0.9037 - val_loss: 0.4419 - val_accuracy: 0.9001 - 824ms/epoch - 8ms/step
```

```
str(history)
```

```
## List of 2
## $ params :List of 3
## ..$ verbose: int 2
## ..$ epochs : int 20
## ..$ steps : int 98
## $ metrics:List of 4
## ..$ loss      : num [1:20] 0.315 0.315 0.322 0.322 0.328 ...
## ..$ accuracy  : num [1:20] 0.915 0.914 0.912 0.911 0.912 ...
## ..$ val_loss   : num [1:20] 0.322 0.378 0.363 0.389 0.341 ...
## ..$ val_accuracy: num [1:20] 0.914 0.9 0.899 0.896 0.912 ...
## - attr(*, "class")= chr "keras_training_history"
```

```
plot(history)
```



```
library(keras)
model <- keras_model_sequential() %>%
layer_conv_2d(filters = 32, kernel_size = c(3, 3), activation = "relu",
input_shape = c(28, 28, 1)) %>%
layer_max_pooling_2d(pool_size = c(2, 2)) %>%
layer_conv_2d(filters = 64, kernel_size = c(3, 3), activation = "relu") %>%
layer_max_pooling_2d(pool_size = c(2, 2)) %>%
layer_conv_2d(filters = 64, kernel_size = c(3, 3), activation = "relu") %>%
layer_flatten() %>%
layer_dense(units = 64, activation = "relu") %>%
layer_dense(units = 10, activation = "softmax")
model
```

```
## Model: "sequential_2"
```

```
##
```

## Layer (type)	Output Shape	Param #
## conv2d_2 (Conv2D)	(None, 26, 26, 32)	320
## max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	0
## conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
## max_pooling2d (MaxPooling2D)	(None, 5, 5, 64)	0
## conv2d (Conv2D)	(None, 3, 3, 64)	36928
## flatten (Flatten)	(None, 576)	0
## dense_5 (Dense)	(None, 64)	36928
## dense_4 (Dense)	(None, 10)	650

```
##
```

Total params: 93322 (364.54 KB)
Trainable params: 93322 (364.54 KB)
Non-trainable params: 0 (0.00 Byte)

```
##
```

```

library(keras)
model_1 <- keras_model_sequential() %>%
  layer_conv_2d(filters = 32, kernel_size = c(5, 5), activation = "relu", input_shape = c(28, 28, 1), padding = "same") %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_conv_2d(filters = 64, kernel_size = c(5, 5), activation = "relu", padding = "same") %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_conv_2d(filters = 64, kernel_size = c(5, 5), activation = "relu", padding = "same") %>%
  layer_flatten() %>%
  layer_dense(units = 64, activation = "relu") %>%
  layer_dense(units = 10, activation = "softmax")

model_1

```

```

## Model: "sequential_3"
##
## Layer (type)                Output Shape                Param #
## =====
## conv2d_5 (Conv2D)           (None, 28, 28, 32)         832
## max_pooling2d_3 (MaxPooling2D) (None, 14, 14, 32)         0
## conv2d_4 (Conv2D)           (None, 14, 14, 64)         51264
## max_pooling2d_2 (MaxPooling2D) (None, 7, 7, 64)           0
## conv2d_3 (Conv2D)           (None, 7, 7, 64)           102464
## flatten_1 (Flatten)         (None, 3136)                0
## dense_7 (Dense)             (None, 64)                  200768
## dense_6 (Dense)             (None, 10)                   650
## =====
## Total params: 355978 (1.36 MB)
## Trainable params: 355978 (1.36 MB)
## Non-trainable params: 0 (0.00 Byte)
##

```

Here we can see the kernel size increases from (3,3) to (5,5). The Param # in result table increases a lot. The same situation happens in total params and trainable params.

```
mnist <- dataset_mnist()
c(c(train_images, train_labels), c(test_images, test_labels)) %<-% mnist
train_images <- array_reshape(train_images, c(60000, 28, 28, 1))
train_images <- train_images / 255
test_images <- array_reshape(test_images, c(10000, 28, 28, 1))
test_images <- test_images / 255
train_labels <- to_categorical(train_labels)
test_labels <- to_categorical(test_labels)
model %>% compile(
  optimizer = "rmsprop",
  loss = "categorical_crossentropy",
  metrics = c("accuracy")
)
model %>% fit(
  train_images, train_labels,
  epochs = 5, batch_size=64
)
```

```
## Epoch 1/5
## 938/938 - 10s - loss: 0.1865 - accuracy: 0.9406 - 10s/epoch - 11ms/step
## Epoch 2/5
## 938/938 - 9s - loss: 0.0557 - accuracy: 0.9828 - 9s/epoch - 10ms/step
## Epoch 3/5
## 938/938 - 9s - loss: 0.0417 - accuracy: 0.9872 - 9s/epoch - 10ms/step
## Epoch 4/5
## 938/938 - 9s - loss: 0.0330 - accuracy: 0.9902 - 9s/epoch - 10ms/step
## Epoch 5/5
## 938/938 - 9s - loss: 0.0278 - accuracy: 0.9923 - 9s/epoch - 10ms/step
```

```
results <- model %>% evaluate(test_images, test_labels)
```

```
## 313/313 - 2s - loss: 0.0351 - accuracy: 0.9896 - 2s/epoch - 6ms/step
```

```
results
```

```
##      loss    accuracy
## 0.03505512 0.98960000
```